

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import random
from google.colab import files
```

```
uploaded = files.upload()
```



Choose Files

 delhi_elect...emapnd.csv

delhi_electricity_demapnd.csv

(text/csv) - 2687390 bytes, last modified: n/a - 100% done

Saving delhi_electricity_demapnd.csv to delhi_electricity_demapnd.csv


```
data = pd.read_csv('delhi_electricity_demapnd.csv')
```

```
data.shape
```



(43848, 9)

```
data.head(-10)
```



	Timestamp	hour	dayofweek	month	year	dayofyear	Temperature	Humi
0	01-Jan-20	0.0	2.0	1.0	2020.0	1.0	3.000000	61.28
1	01-Jan-20	1.0	2.0	1.0	2020.0	1.0	3.000000	52.87
2	01-Jan-20	2.0	2.0	1.0	2020.0	1.0	4.244482	36.34
3	01-Jan-20	3.0	2.0	1.0	2020.0	1.0	3.000000	72.62
4	01-Jan-20	4.0	2.0	1.0	2020.0	1.0	3.881208	90.58
...
43833	31-Dec-24	9.0	1.0	12.0	2024.0	366.0	16.896276	72.38
43834	31-Dec-24	10.0	1.0	12.0	2024.0	366.0	21.984042	77.27
43835	31-Dec-24	11.0	1.0	12.0	2024.0	366.0	20.659219	69.22
43836	31-Dec-24	12.0	1.0	12.0	2024.0	366.0	19.112534	52.68
43837	31-Dec-24	13.0	1.0	12.0	2024.0	366.0	17.862808	64.74

43838 rows x 9 columns

Next steps:

[Generate code with data](#)

 [View recommended plots](#)

[New interactive sheet](#)

```
data.info()
```

```
↗ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 43848 entries, 0 to 43847
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Timestamp       43848 non-null  object
1   hour            43837 non-null  float64
2   dayofweek       43839 non-null  float64
3   month           43840 non-null  float64
4   year            43843 non-null  float64
5   dayofyear       43843 non-null  float64
6   Temperature     43841 non-null  float64
7   Humidity        43838 non-null  float64
8   Demand          43841 non-null  float64
dtypes: float64(8), object(1)
memory usage: 3.0+ MB
```

the **Timestamp** column was initially a an object data type and below we are converting it to *datetime* format.

```
data['Timestamp'] = pd.to_datetime(data['Timestamp'])
data.info()
```

```
↗ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 43848 entries, 0 to 43847
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Timestamp       43848 non-null  datetime64[ns]
1   hour            43837 non-null  float64
2   dayofweek       43839 non-null  float64
3   month           43840 non-null  float64
4   year            43843 non-null  float64
5   dayofyear       43843 non-null  float64
6   Temperature     43841 non-null  float64
7   Humidity        43838 non-null  float64
8   Demand          43841 non-null  float64
dtypes: datetime64[ns](1), float64(8)
memory usage: 3.0 MB
/tmp/ipython-input-7-3315307892.py:1: UserWarning: Could not infer format,
  data['Timestamp'] = pd.to_datetime(data['Timestamp'])
```

```
data = data.set_index('Timestamp')
data
```

	hour	dayofweek	month	year	dayofyear	Temperature	Humidity	
Timestamp								
2020-01-01	0.0	2.0	1.0	2020.0	1.0	3.000000	61.288951	2
2020-01-01	1.0	2.0	1.0	2020.0	1.0	3.000000	52.873702	2
2020-01-01	2.0	2.0	1.0	2020.0	1.0	4.244482	36.341783	2
2020-01-01	3.0	2.0	1.0	2020.0	1.0	3.000000	72.629378	2
2020-01-01	4.0	2.0	1.0	2020.0	1.0	3.881208	90.582444	2
...
2024-12-31	19.0	1.0	12.0	2024.0	366.0	3.956838	43.287161	4
2024-12-31	20.0	1.0	12.0	2024.0	366.0	3.118824	51.705756	4
2024-12-31	21.0	1.0	12.0	2024.0	366.0	3.000000	40.565916	4
2024-12-31	22.0	1.0	12.0	2024.0	366.0	3.000000	51.998107	3
2024-12-31	23.0	1.0	12.0	2024.0	366.0	6.037472	59.931925	3

43848 rows × 8 columns

Next steps:

Generate code with data

View recommended plots

New interactive sheet

```
data[['Temperature', 'Humidity', 'Demand']].describe()
```

	Temperature	Humidity	Demand
count	43841.000000	43838.000000	43841.000000
mean	25.067788	59.903007	5000.790976
std	12.821725	18.342604	1412.527409
min	3.000000	20.000000	1611.954020
25%	15.210186	46.241224	4015.668472
50%	25.003212	59.986720	5013.053367
75%	34.740971	73.796820	6000.803082
max	50.000000	95.000000	11910.705100

✓ Dealing with Missing Data

```
print(data.isnull().sum())
```

```
↔ hour          11  
   dayofweek     9  
   month         8  
   year          5  
   dayofyear     5  
   Temperature   7  
   Humidity      10  
   Demand        7  
   dtype: int64
```

```
data[data.isna().any(axis=1)]
```

```
↔
```

	hour	dayofweek	month	year	dayofyear	Temperature	Humidity
Timestamp							
2020-04-30	NaN	3.0	4.0	2020.0	121.0	21.820261	41.353675
2020-07-21	NaN	1.0	7.0	2020.0	203.0	36.555833	62.779665
2021-02-20	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2021-02-20	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2021-02-20	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2021-02-20	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2021-02-27	NaN	5.0	2.0	2021.0	58.0	24.001003	53.303268
2021-12-13	23.0	0.0	12.0	2021.0	NaN	3.000000	75.457130
2021-12-14	NaN	1.0	12.0	2021.0	348.0	5.061319	47.215825
2021-12-14	14.0	1.0	12.0	2021.0	348.0	15.403995	74.011557
2022-06-17	13.0	4.0	6.0	2022.0	168.0	NaN	43.171391
2022-06-17	NaN	4.0	6.0	2022.0	168.0	42.924693	43.645711
2022-06-17	17.0	4.0	NaN	2022.0	168.0	37.115634	37.554842
2022-06-17	18.0	4.0	6.0	2022.0	168.0	39.886471	NaN
2022-06-17	21.0	NaN	6.0	2022.0	168.0	28.471494	54.964101
2022-06-18	5.0	5.0	6.0	2022.0	169.0	38.349174	58.285814
2023-08-30	NaN	2.0	8.0	2023.0	242.0	25.707715	69.868663
2023-08-31	0.0	NaN	8.0	2023.0	243.0	22.705122	81.691817

2023-08-31	5.0	3.0	NaN	2023.0	243.0	29.081792	81.270824
2023-08-31	11.0	3.0	8.0	NaN	243.0	44.270455	75.563954
2023-08-31	14.0	3.0	8.0	2023.0	243.0	NaN	NaN
2024-01-27	19.0	5.0	1.0	2024.0	27.0	14.351975	NaN
2024-01-28	2.0	NaN	1.0	2024.0	28.0	3.000000	68.899943
2024-01-28	NaN	6.0	1.0	2024.0	28.0	9.081186	78.163106
2024-09-24	19.0	NaN	9.0	2024.0	268.0	25.632052	74.123843
2024-09-24	21.0	1.0	9.0	2024.0	268.0	19.380978	NaN
2024-09-24	22.0	1.0	9.0	2024.0	268.0	NaN	66.618690
2024-09-25	7.0	2.0	NaN	2024.0	269.0	31.190258	95.000000
2024-12-18	11.0	2.0	12.0	2024.0	353.0	23.075591	NaN
2024-12-18	12.0	2.0	12.0	2024.0	353.0	20.012449	NaN
2024-12-18	21.0	NaN	NaN	2024.0	353.0	3.000000	64.041109

```
data[data.isna().all(axis=1)]
##this shows the subset of the data where all the values in each of the columns
```


↔

	hour	dayofweek	month	year	dayofyear	Temperature	Humidity	De
Timestamp								
2021-02-20	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
2021-02-20	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
2021-02-20	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
2021-02-20	NaN	NaN	NaN	NaN	NaN	NaN	NaN	

As there is no scope of trying to work our way around with the data where no entries are present, it would be a wise decision to drop these rows of data.

Below we are dropping all the rows which had no data in them at all!

```
data = data.dropna(how='all')
data
```



	hour	dayofweek	month	year	dayofyear	Temperature	Humidity	
Timestamp								
2020-01-01	0.0	2.0	1.0	2020.0	1.0	3.000000	61.288951	2
2020-01-01	1.0	2.0	1.0	2020.0	1.0	3.000000	52.873702	2
2020-01-01	2.0	2.0	1.0	2020.0	1.0	4.244482	36.341783	2
2020-01-01	3.0	2.0	1.0	2020.0	1.0	3.000000	72.629378	2
2020-01-01	4.0	2.0	1.0	2020.0	1.0	3.881208	90.582444	2
...
2024-12-31	19.0	1.0	12.0	2024.0	366.0	3.956838	43.287161	4
2024-12-31	20.0	1.0	12.0	2024.0	366.0	3.118824	51.705756	4
2024-12-31	21.0	1.0	12.0	2024.0	366.0	3.000000	40.565916	4
2024-12-31	22.0	1.0	12.0	2024.0	366.0	3.000000	51.998107	3
2024-12-31	23.0	1.0	12.0	2024.0	366.0	6.037472	59.931925	3

43844 rows x 8 columns

Next steps:


[Generate code with data](#)

[View recommended plots](#)

[New interactive sheet](#)

Now to fill in the remaining empty values we can use bfill and ffill

```
data[['hour', 'dayofweek', 'month', 'year', 'dayofyear']] = data[['hour', 'dayc
```



```
/tmp/ipython-input-14-3402763247.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs>
data[['hour', 'dayofweek', 'month', 'year', 'dayofyear']] = data[['hour',

```
data[['Temperature','Humidity']] = data[['Temperature','Humidity']].bfill()
```

↵ /tmp/ipython-input-15-2229112482.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs>
data[['Temperature','Humidity']] = data[['Temperature','Humidity']].bfill

```
data['Demand'] = data['Demand'].interpolate(method='time')
```

↵ /tmp/ipython-input-16-3652729601.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs>
data['Demand'] = data['Demand'].interpolate(method='time')

```
data.isnull().sum()
```


↵

	0
<hr/>	
hour	0
dayofweek	0
month	0
year	0
dayofyear	0
Temperature	0
Humidity	0
Demand	0

dtype: int64

✓ Feature Engineering

data



	hour	dayofweek	month	year	dayofyear	Temperature	Humidity	
Timestamp								
2020-01-01	0.0	2.0	1.0	2020.0	1.0	3.000000	61.288951	2
2020-01-01	1.0	2.0	1.0	2020.0	1.0	3.000000	52.873702	2
2020-01-01	2.0	2.0	1.0	2020.0	1.0	4.244482	36.341783	2
2020-01-01	3.0	2.0	1.0	2020.0	1.0	3.000000	72.629378	2
2020-01-01	4.0	2.0	1.0	2020.0	1.0	3.881208	90.582444	2
...
2024-12-31	19.0	1.0	12.0	2024.0	366.0	3.956838	43.287161	4
2024-12-31	20.0	1.0	12.0	2024.0	366.0	3.118824	51.705756	4
2024-12-31	21.0	1.0	12.0	2024.0	366.0	3.000000	40.565916	4
2024-12-31	22.0	1.0	12.0	2024.0	366.0	3.000000	51.998107	3
2024-12-31	23.0	1.0	12.0	2024.0	366.0	6.037472	59.931925	3

43844 rows x 8 columns


Next steps:

[Generate code with data](#)

[View recommended plots](#)

[New interactive sheet](#)


```
data.insert(5, 'quarter', data.index.quarter)
data
```



	hour	dayofweek	month	year	dayofyear	quarter	Temperature	Hu
Timestamp								
2020-01-01	0.0	2.0	1.0	2020.0	1.0	1	3.000000	61
2020-01-01	1.0	2.0	1.0	2020.0	1.0	1	3.000000	52
2020-01-01	2.0	2.0	1.0	2020.0	1.0	1	4.244482	36
2020-01-01	3.0	2.0	1.0	2020.0	1.0	1	3.000000	72
2020-01-01	4.0	2.0	1.0	2020.0	1.0	1	3.881208	90
...
2024-12-31	19.0	1.0	12.0	2024.0	366.0	4	3.956838	43
2024-12-31	20.0	1.0	12.0	2024.0	366.0	4	3.118824	51
2024-12-31	21.0	1.0	12.0	2024.0	366.0	4	3.000000	40
2024-12-31	22.0	1.0	12.0	2024.0	366.0	4	3.000000	51
2024-12-31	23.0	1.0	12.0	2024.0	366.0	4	6.037472	59

43844 rows × 9 columns


Next steps:

[Generate code with data](#)

[View recommended plots](#)

[New interactive sheet](#)

```
data[['hour', 'dayofweek', 'month', 'year', 'dayofyear']] = data[['hour', 'dayc
```



/tmp/ipython-input-20-3222519661.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs>
data[['hour', 'dayofweek', 'month', 'year', 'dayofyear']] = data[['hour',

data.info()

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 43844 entries, 2020-01-01 to 2024-12-31
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   hour                   43844 non-null  int64
1   dayofweek              43844 non-null  int64
2   month                  43844 non-null  int64
3   year                   43844 non-null  int64
4   dayofyear              43844 non-null  int64
5   quarter                43844 non-null  int32
6   Temperature            43844 non-null  float64
7   Humidity               43844 non-null  float64
8   Demand                 43844 non-null  float64
dtypes: float64(3), int32(1), int64(5)
memory usage: 3.2 MB
```

data.insert(5, 'weekofyear', data.index.isocalendar().week.astype(int))


data

	hour	dayofweek	month	year	dayofyear	weekofyear	quarter	Temp
Timestamp								
2020-01-01	0	2	1	2020	1	1	1	
2020-01-01	1	2	1	2020	1	1	1	
2020-01-01	2	2	1	2020	1	1	1	
2020-01-01	3	2	1	2020	1	1	1	
2020-01-01	4	2	1	2020	1	1	1	
...	
2024-12-31	19	1	12	2024	366	1	4	
2024-12-31	20	1	12	2024	366	1	4	
2024-12-31	21	1	12	2024	366	1	4	
2024-12-31	22	1	12	2024	366	1	4	
2024-12-31	23	1	12	2024	366	1	4	

43844 rows x 10 columns

```
data.insert(7, 'is_weekend' , data.index.dayofweek.isin([5,6]).astype(int))
```

data



	hour	dayofweek	month	year	dayofyear	weekofyear	quarter	is_w
Timestamp								
2020-01-01	0	2	1	2020	1	1	1	
2020-01-01	1	2	1	2020	1	1	1	
2020-01-01	2	2	1	2020	1	1	1	
2020-01-01	3	2	1	2020	1	1	1	
2020-01-01	4	2	1	2020	1	1	1	
...
2024-12-31	19	1	12	2024	366	1	4	
2024-12-31	20	1	12	2024	366	1	4	
2024-12-31	21	1	12	2024	366	1	4	
2024-12-31	22	1	12	2024	366	1	4	
2024-12-31	23	1	12	2024	366	1	4	

43844 rows x 11 columns

Next steps:

[Generate code with data](#)


[View recommended plots](#)

[New interactive sheet](#)

Also we will check if there was any public holiday during the timestamp of the dataset

```
import holidays
```

```
data['holidays'] = holidays.IN(years = data.year)
```



/tmp/ipython-input-27-2470672226.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs>
data['holidays'] = holidays.IN(years = data.year)

```
data.holidays.value_counts()
```



	count
holidays	
Republic Day	120
Maha Shivaratri	120
Mahavir Jayanti	120
Good Friday	120
Buddha Purnima	120
Eid al-Fitr	120
Eid al-Adha	120
Janmashtami	120
Independence Day	120
Ashura	120
Gandhi Jayanti	120
Dussehra	120
Prophet's Birthday	120
Diwali	120
Guru Nanak Jayanti	120
Christmas	120

dtype: int64

```
data = data.drop(columns = ['holidays'])
```

```
data['Demand_lag_24hr'] = data['Demand'].shift(24)
```


```
data['Demand_lag_week168hr'] = data['Demand'].shift(168)
```

✓ Rolling Meand and Rolling Standard Deviation

```
data['Demand_rolling_mean_24hr'] = data['Demand'].rolling(window=24).mean()
```

```
data['Demand_rolling_std_24hr'] = data['Demand'].rolling(window=24).std()
```

```
data.head(30)
```



	hour	dayofweek	month	year	dayofyear	weekofyear	quarter	is_w
Timestamp								
2020-01-01	0	2	1	2020	1	1	1	
2020-01-01	1	2	1	2020	1	1	1	
2020-01-01	2	2	1	2020	1	1	1	
2020-01-01	3	2	1	2020	1	1	1	
2020-01-01	4	2	1	2020	1	1	1	
2020-01-01	5	2	1	2020	1	1	1	
2020-01-01	6	2	1	2020	1	1	1	
2020-01-01	7	2	1	2020	1	1	1	
2020-01-01	8	2	1	2020	1	1	1	
2020-01-01	9	2	1	2020	1	1	1	
2020-01-01	10	2	1	2020	1	1	1	
2020-01-01	11	2	1	2020	1	1	1	
2020-01-01	12	2	1	2020	1	1	1	
2020-01-01	13	2	1	2020	1	1	1	
2020-01-01	14	2	1	2020	1	1	1	
2020-01-01	15	2	1	2020	1	1	1	
2020-01-01	16	2	1	2020	1	1	1	
2020-01-01	17	2	1	2020	1	1	1	
2020-01-01	18	2	1	2020	1	1	1	
2020-01-01	19	2	1	2020	1	1	1	
2020-01-01	20	2	1	2020	1	1	1	
2020-01-01	21	2	1	2020	1	1	1	
2020-01-01	22	2	1	2020	1	1	1	
2020-01-01	23	2	1	2020	1	1	1	
2020-01-02	0	3	1	2020	2	1	1	
2020-01-02	1	3	1	2020	2	1	1	

2020-01-02	2	3	1	2020	2	1	1
2020-01-02	3	3	1	2020	2	1	1
2020-01-02	4	3	1	2020	2	1	1
2020-01-02	5	3	1	2020	2	1	1

Next steps:


Generate code with data

View recommended plots

New interactive sheet

```
##now dropping all the rows with null entries
data = data.dropna()
```

data



	hour	dayofweek	month	year	dayofyear	weekofyear	quarter	is_w
Timestamp								
2020-01-08	0	2	1	2020	8	2	1	
2020-01-08	1	2	1	2020	8	2	1	
2020-01-08	2	2	1	2020	8	2	1	
2020-01-08	3	2	1	2020	8	2	1	
2020-01-08	4	2	1	2020	8	2	1	
...	
2024-12-31	19	1	12	2024	366	1	4	
2024-12-31	20	1	12	2024	366	1	4	
2024-12-31	21	1	12	2024	366	1	4	
2024-12-31	22	1	12	2024	366	1	4	
2024-12-31	23	1	12	2024	366	1	4	

43676 rows x 15 columns

Next steps:

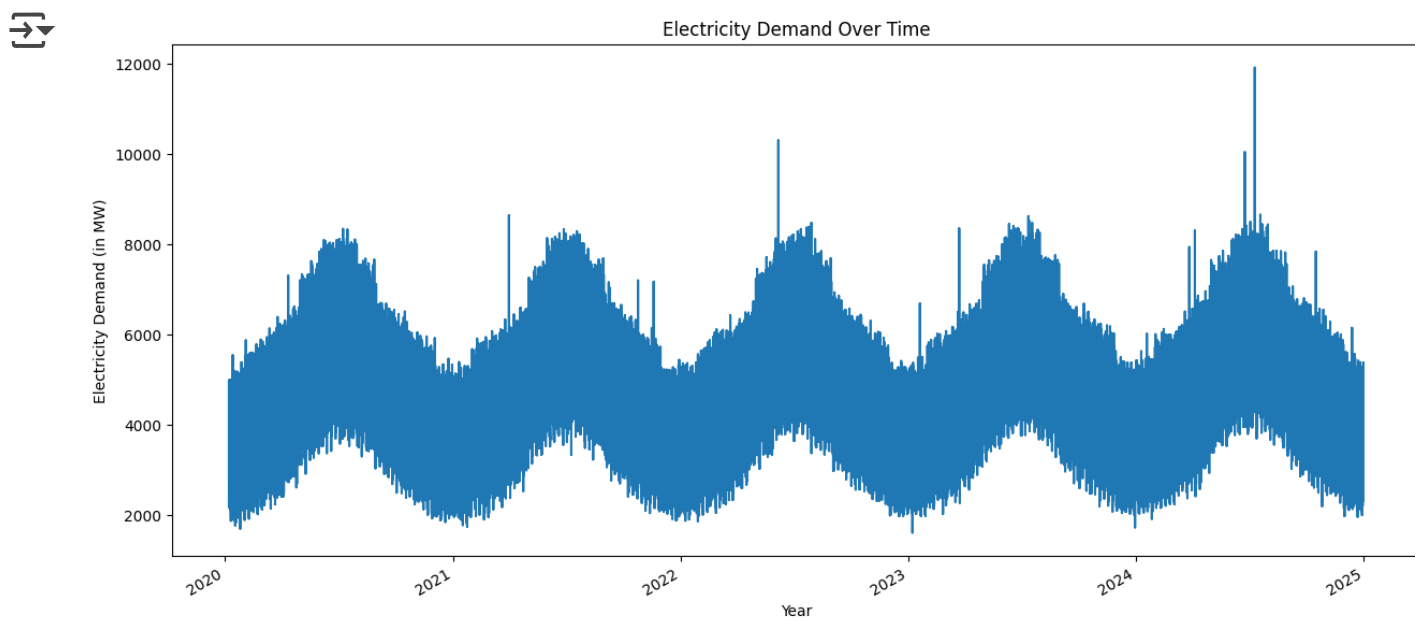
Generate code with data

View recommended plots


New interactive sheet

Visualization

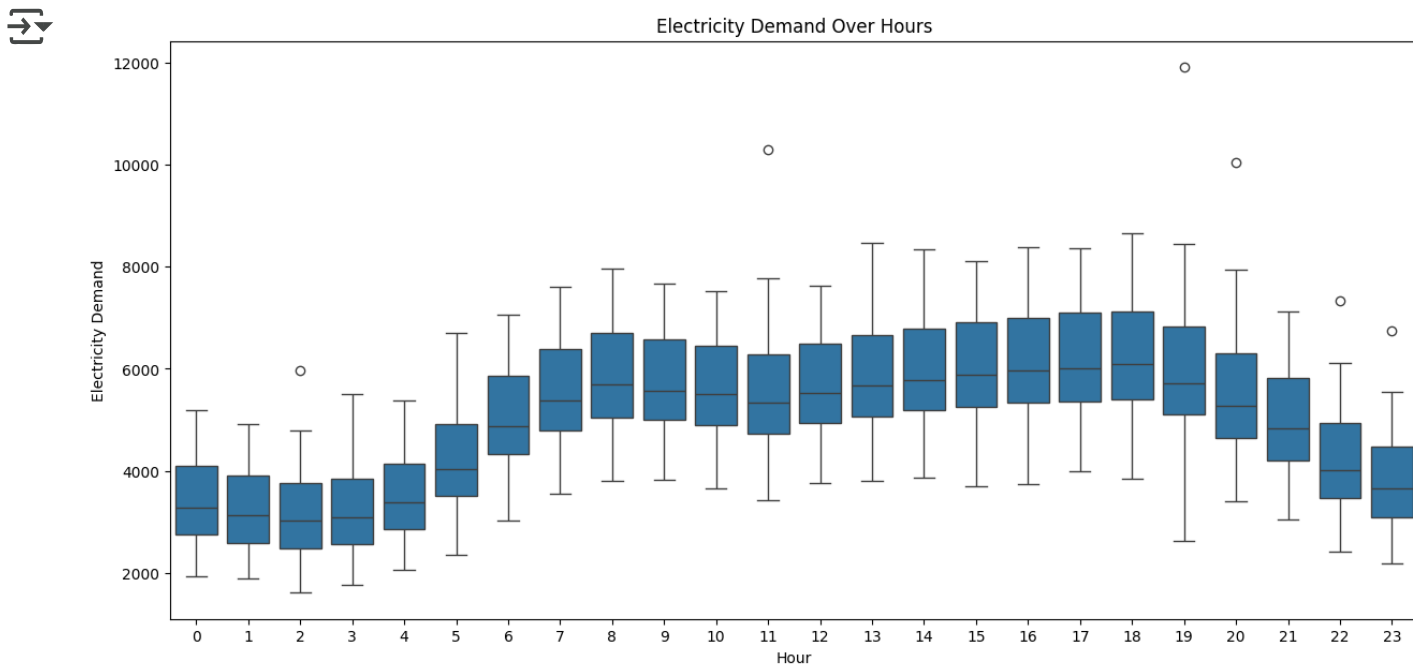
```
data['Demand'].plot(title = 'Electricity Demand Over Time', figsize = (15,7))
plt.xlabel('Year')
plt.ylabel('Electricity Demand (in MW)')
plt.show()
```



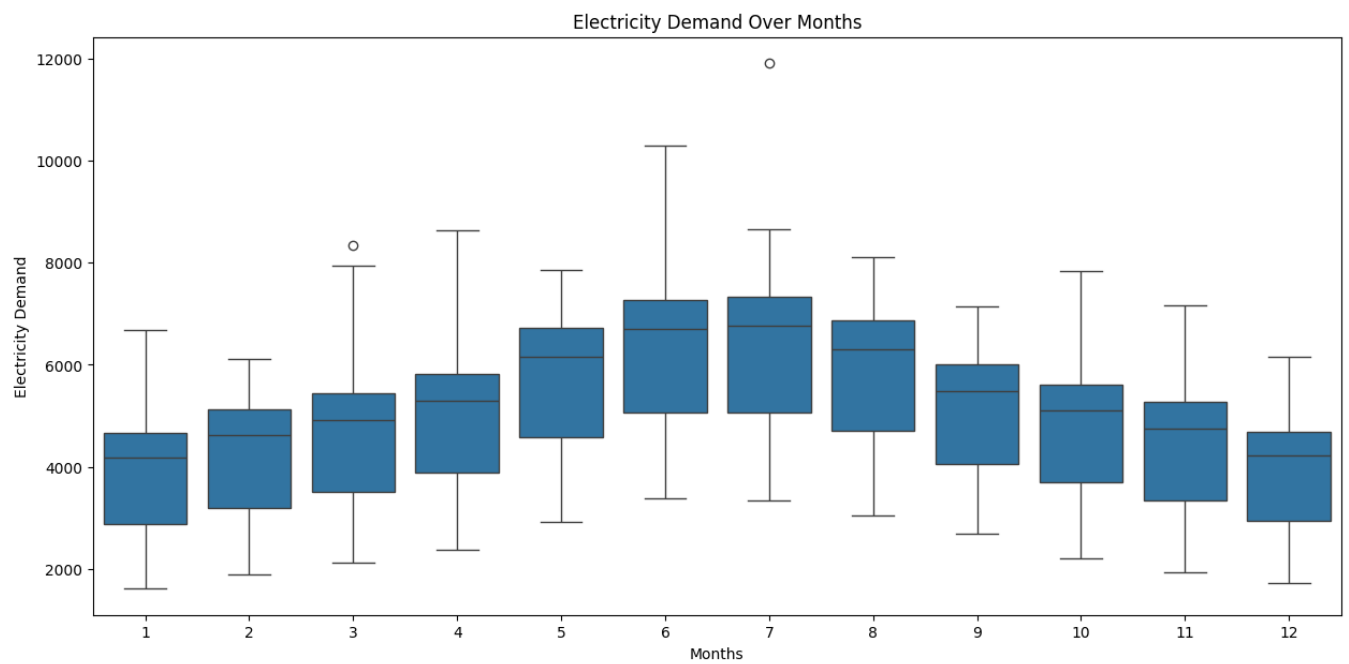
```
sns.barplot(x=data.is_weekend, y=data.Demand)
plt.title('Electricity Demand Over Weekends')
plt.xticks([0, 1], ['Weekday', 'Weekend'])
plt.xlabel('Phase of Week')
plt.ylabel('Electricity Demand')
plt.show()
```



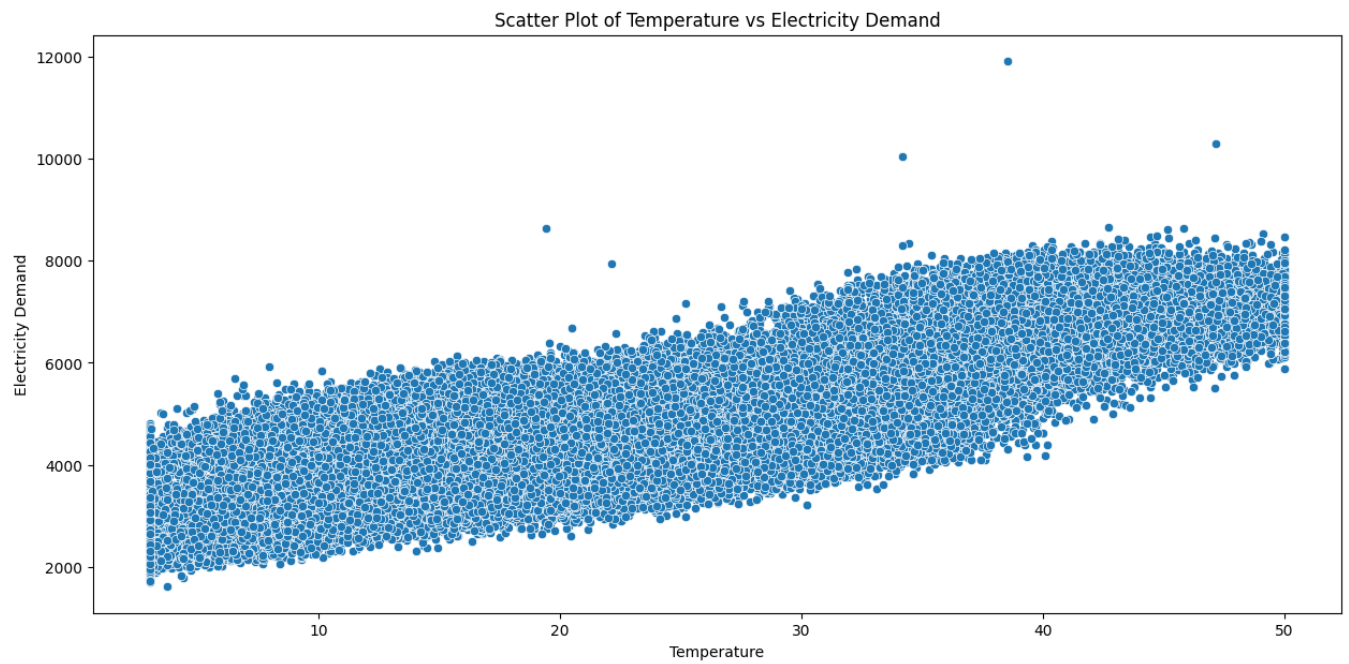

```
plt.figure(figsize=(15,7))
sns.boxplot(x=data.hour, y=data.Demand)
plt.title('Electricity Demand Over Hours')
plt.xlabel('Hour')
plt.ylabel('Electricity Demand')
plt.show()
```



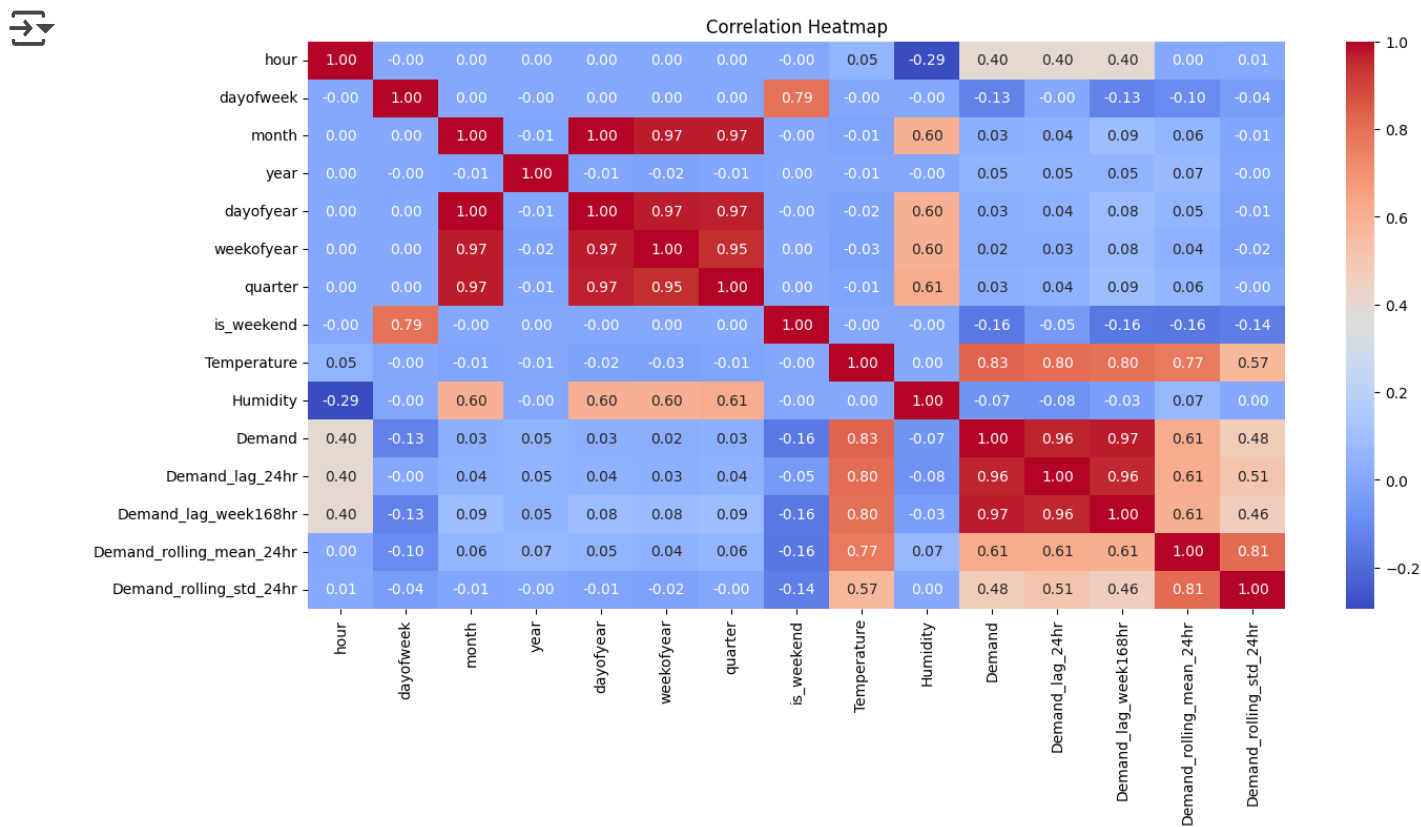
```
plt.figure(figsize=(15,7))
sns.boxplot(x=data.month, y=data.Demand)
plt.title('Electricity Demand Over Months')
plt.xlabel('Months')
plt.ylabel('Electricity Demand')
plt.show()
```



```
plt.figure(figsize=(15,7))
sns.scatterplot(x=data.Temperature, y=data.Demand)
plt.title('Scatter Plot of Temperature vs Electricity Demand')
plt.xlabel('Temperature')
plt.ylabel('Electricity Demand')
plt.show()
```



```
plt.figure(figsize=(15,7))
sns.heatmap(data.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap')
plt.show()
```



Correlation heatmaps shows how one column's values change w.r.t other column according to their correlation points (CP).


- when CP = 1, it means to columns are directly correlated
- when CP = -1, it means to columns are inversely correlated

✓ **Model Building**

```
# assigning dependen variable
y = data.Demand
```

```
# dependent variables consists of all the columns(parameters) except the Demanc
x = data.drop(columns = ['Demand'], axis=1)
```

y




Demand	
Timestamp	
2020-01-08	2363.060115
2020-01-08	2282.558766
2020-01-08	2193.324174
2020-01-08	2208.724679
2020-01-08	2402.611018
...	...
2024-12-31	4689.693109
2024-12-31	4331.249224
2024-12-31	4015.979957
2024-12-31	3353.241682
2024-12-31	3219.023339

43676 rows × 1 columns

dtype: float64

x



	hour	dayofweek	month	year	dayofyear	weekofyear	quarter	is_w
Timestamp								
2020-01-08	0	2	1	2020	8	2	1	
2020-01-08	1	2	1	2020	8	2	1	
2020-01-08	2	2	1	2020	8	2	1	
2020-01-08	3	2	1	2020	8	2	1	
2020-01-08	4	2	1	2020	8	2	1	
...
2024-12-31	19	1	12	2024	366	1	4	
2024-12-31	20	1	12	2024	366	1	4	
2024-12-31	21	1	12	2024	366	1	4	
2024-12-31	22	1	12	2024	366	1	4	
2024-12-31	23	1	12	2024	366	1	4	

43676 rows x 14 columns

Next steps:

[Generate code with x](#)

 [View recommended plots](#)

[New interactive sheet](#)

Splitting the data into 80 - 20 for training and testing

```
x_train = x.loc[:'2023-12-31']

y_train = y.loc[:'2023-12-31']

x_test = x.loc['2024-01-01':]

y_test = y.loc['2024-01-01':]
```

✓ XGBoost Model

As it can handle non-linear data easily for timeseries forecasting

```

from xgboost import XGBRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.model_selection import TimeSeriesSplit

```

Training the model

```

model_xgb = XGBRegressor(n_estimators=1000,
                          learning_rate=0.01,
                          early_stopping_rounds=50,
                          random_state=42,
                          objective='reg:squarederror')

```

In short:

The XGBRegressor initialization creates a model that will train for up to 1000 boosting rounds (trees) but will stop early if performance on a validation set doesn't improve for 50 consecutive rounds.


Each tree's contribution is scaled down by a learning rate of 0.01 to prevent overfitting.

This model is designed to minimize the squared error for regression tasks.

```

model_xgb.fit(x_train, y_train, eval_set=[(x_test, y_test)], verbose=False)

```



```

XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=50,
              enable_categorical=False, eval_metric=None, feature_types=None,
              feature_weights=None, gamma=None, grow_policy=None,
              importance_type=None, interaction_constraints=None,
              learning_rate=0.01, max_bin=None, max_cat_threshold=None,
              max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
              max_leaves=None, min_child_weight=None, missing=nan,
              monotone_constraints=None, multi_strategy=None, n_estimators=1000,
              n_jobs=None, num_parallel_tree=None, ...)

```

```

# making predictions
y_pred = model_xgb.predict(x_test)

# evaluating its findings
root_mse = np.sqrt(mean_squared_error(y_test, y_pred))
mae = mean_absolute_error(y_test, y_pred)

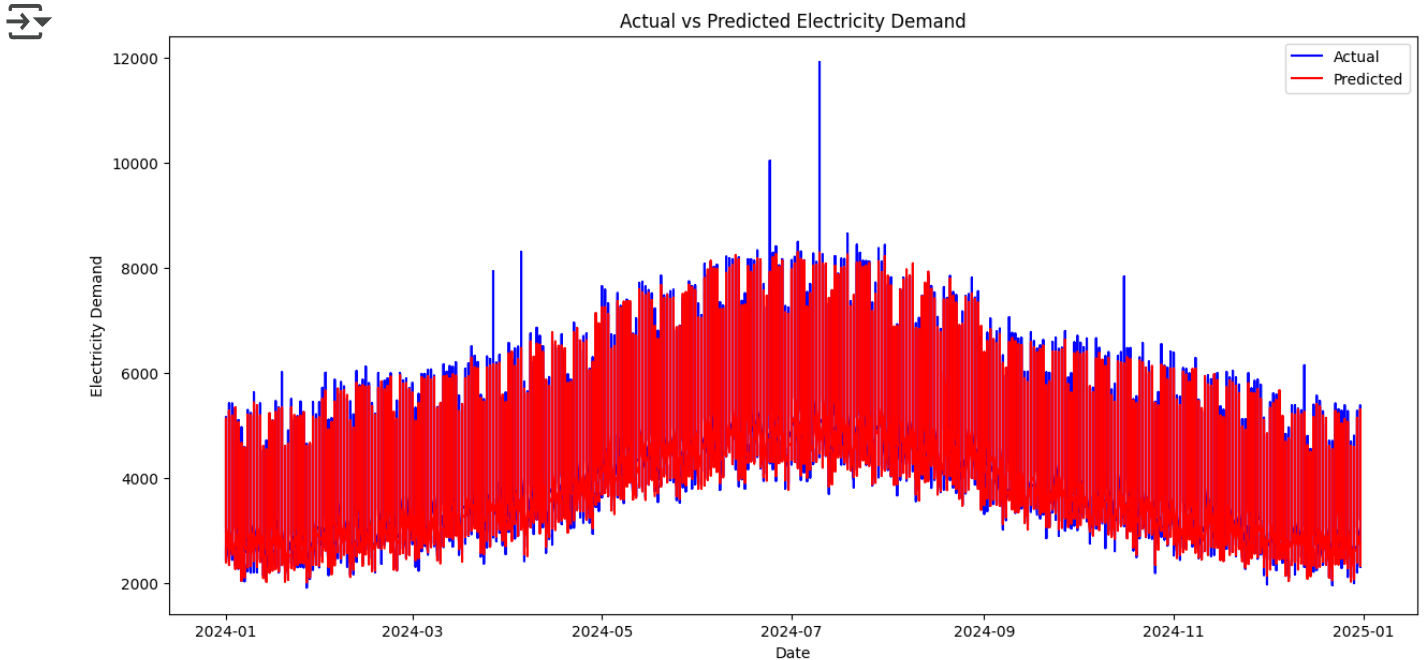
```

```
print("Root Mean Squared Error:", root_mse)
print("Mean Absolute Error:", mae)
```

```
↔ Root Mean Squared Error: 175.22716387271916
   Mean Absolute Error: 123.47612356015286
```

✓ Visualising the Model Predictions


```
plt.figure(figsize=(15,7))
plt.plot(y_test.index, y_test.values, label='Actual', color='blue')
plt.plot(y_test.index, y_pred, label='Predicted', color='red')
plt.title('Actual vs Predicted Electricity Demand')
plt.xlabel('Date')
plt.ylabel('Electricity Demand')
plt.legend()
plt.show()
```



```
# saving the model
import joblib
joblib.dump(model_xgb, 'xgb_model.pkl')
```

```
['xgb_model.pkl']
```

```
files.download('xgb_model.pkl')
```



We can use this .pkl file to instantiate the model we just created.

```
# to load the pkl file
model_new = joblib.load('xgb_model.pkl')
```

```
# verify the model
model_new
```



```
▼ XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
             colsample_bylevel=None, colsample_bynode=None,
             colsample_bytree=None, device=None, early_stopping_rounds=50,
             enable_categorical=False, eval_metric=None, feature_types=None,
             feature_weights=None, gamma=None, grow_policy=None,
             importance_type=None, interaction_constraints=None,
             learning_rate=0.01, max_bin=None, max_cat_threshold=None,
             max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
             max_leaves=None, min_child_weight=None, missing=nan,
             monotone_constraints=None, multi_strategy=None, n_estimators=100,
             n_jobs=None, num_parallel_tree=None, ...)
```